

Assignment 3 report

github repo:

<https://github.com/yty20/CS6650-project>

Database Designs

I uses DynamoDB and there is only one table named "skierActivities" with the following attributes:

- skierID (String): partition key
- date (String): sort key
- lifts (Map of String(liftID) and Integer(number of lifts))
- resortID (String)
- verticalTotals (integer): all liftID in lifts times 10.

And I also add a GSI with resortID as partition key and date as sort key to help query by resortID and date.

Deployment Topology

The deployment topology is as follows:

client -> EC2 with Servlet -> RMQ -> consumer * 4 -> DynamoDB.


the client use 100 threads to simulate the skiers' activities. Servlet collects the data from the skiers and send it to RMQ. There are 4 instances of consumer to process the data from RMQ and store it in DynamoDB, so that the consumer rate can keep up with the incoming data and keep the queue in healthy status.

Metric Results:

1. RabbitMQ console log with only one consumer:

Overview


Totals

Queued messages [last minute](#) ?

Ready 41,388

Unacked 298

Total 41,686

Message rates [last minute](#) ?

Publish 0.00/s

Publisher confirm 0.00/s

Deliver (manual ack) 858/s

Deliver (auto ack) 0.00/s

Consumer ack 858/s

Redelivered 0.00/s

Get (manual ack) 0.00/s

Get (auto ack) 0.00/s

Get (empty) 0.00/s

Unroutable (return) 0.00/s

Unroutable (drop) 0.00/s

Disk read 819/s

Disk write 0.00/s

Global counts ?

Connections: 401

Channels: 498

Exchanges: 7

Queues: 1

Consumers: 298

Nodes


Name	File descriptors ?	Socket descriptors ?	Erlang processes	Memory ?	Disk space	Uptime	Info	Reset stats	+/-
rabbit@ip-172-31-19-197	439 65536 available	401 58893 available	5156 1048576 available	252 MiB 380 MiB high watermark	3.8 GiB 48 MiB low watermark	2m 39s	basic disc 1 rss	This node All nodes	

Churn statistics

2. RabbitMQ console log with 4 consumers:

Overview


Totals

Queued messages [last minute](#) ?

Ready 0

Unacked 516

Total 516

Message rates [last minute](#) ?

Publish 1,409/s

Publisher confirm 0.00/s

Deliver (manual ack) 1,395/s

Deliver (auto ack) 0.00/s

Consumer ack 1,450/s

Redelivered 0.00/s

Get (manual ack) 0.00/s

Get (auto ack) 0.00/s

Get (empty) 0.00/s

Unroutable (return) 0.00/s

Unroutable (drop) 0.00/s

Disk read 0.00/s

Disk write 0.00/s

Global counts ?

Connections: 1001

Channels: 954

Exchanges: 7

Queues: 1

Consumers: 754

Nodes

Name	File descriptors ?	Socket descriptors ?	Erlang processes	Memory ?	Disk space	Uptime	Info	Reset stats	+/-
rabbit@ip-172-31-19-197	1039 65536 available	1001 58893 available	11180 1048576 available	360 MiB 380 MiB high watermark	3.8 GiB 48 MiB low watermark	1m 24s	basic disc 1 rss	This node All nodes	

3. client output:

```
Total Successful Requests: 200000
Total Failed Requests: 0
Total Time: 169.603 s
Total throughput: 1179.2244240962718 requests/s
Latency Statistics:
Mean Latency: 84.39164 ms
Median Latency: 59 ms
Min Latency: 31 ms
Max Latency: 4203 ms
99th Percentile Latency: 250 ms

Process finished with exit code 0
```